

CLASS WordCounCREATEt.java

<https://hadoop.apache.org/docs/current/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class IntSumReducer
        extends Reducer<Text,IntWritable,Text,IntWritable> {
        private IntWritable result = new IntWritable();

        public void reduce(Text key, Iterable<IntWritable> values,
            Context context
            ) throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            result.set(sum);
        }
    }
}
```

```

        context.write(key, result);
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

TERMINAL

```

[cloudera@quickstart ~]$ ls
cloudera-manager eclipse Music Templates
cm_api.py enterprise-deployment.json parcels Videos
Desktop express-deployment.json Pictures WordCount.jar
Documents kerberos Processfile.txt WordCountjar.jar
Downloads lib Public workspace
[cloudera@quickstart ~]$ pwd
/home/cloudera
[cloudera@quickstart ~]$ cat > /home/cloudera/devansh.txt
devansh
jay shree ram
bharat mata ke jay
hp boss
om papadiwal
rohit patil
^Z
[1]+ Stopped cat > /home/cloudera/devansh.txt
[cloudera@quickstart ~]$ cat /home/cloudera/devansh.txt
devansh
jay shree ram
bharat mata ke jay
hp boss
om papadiwal
rohit patil
[cloudera@quickstart ~]$ hdfs dfs -ls
[cloudera@quickstart ~]$ hdfs dfs -ls /

```

Found 7 items

```
drwxrwxrwx - hdfs supergroup          0 2017-10-23 09:15 /benchmarks
drwxr-xr-x - hbase supergroup          0 2025-04-22 01:22 /hbase
drwxr-xr-x - cloudera supergroup        0 2025-03-27 23:38 /inputfolder1
drwxr-xr-x - solr solr                  0 2017-10-23 09:18 /solr
drwxrwxrwt - hdfs supergroup          0 2025-03-26 22:24 /tmp
drwxr-xr-x - hdfs supergroup          0 2017-10-23 09:17 /user
drwxr-xr-x - hdfs supergroup          0 2017-10-23 09:17 /var
```

[cloudera@quickstart ~]\$ hdfs dfs -mkdir /input

[cloudera@quickstart ~]\$ hdfs dfs -put /home/cloudera/devansh.txt /input/

[cloudera@quickstart ~]\$ hdfs dfs -ls /

Found 8 items

```
drwxrwxrwx - hdfs supergroup          0 2017-10-23 09:15 /benchmarks
drwxr-xr-x - hbase supergroup          0 2025-04-22 01:22 /hbase
drwxr-xr-x - cloudera supergroup        0 2025-04-22 01:33 /input
drwxr-xr-x - cloudera supergroup        0 2025-03-27 23:38 /inputfolder1
drwxr-xr-x - solr solr                  0 2017-10-23 09:18 /solr
drwxrwxrwt - hdfs supergroup          0 2025-03-26 22:24 /tmp
drwxr-xr-x - hdfs supergroup          0 2017-10-23 09:17 /user
drwxr-xr-x - hdfs supergroup          0 2017-10-23 09:17 /var
```

[cloudera@quickstart ~]\$ hdfs dfs -cat /input/devansh.txt

devansh

jay shree ram

bharat mata ke jay

hp boss

om papadiwal

rohit patil

[cloudera@quickstart ~]\$ hadoop jar /home/cloudera/WordCount.jar WordCount

/input/devansh.txt /out1

25/04/22 01:36:50 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

25/04/22 01:36:51 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.

25/04/22 01:36:52 INFO input.FileInputFormat: Total input paths to process : 1

25/04/22 01:36:52 WARN hdfs.DFSClient: Caught exception

java.lang.InterruptedExecution

at java.lang.Object.wait(Native Method)

at java.lang.Thread.join(Thread.java:1281)

at java.lang.Thread.join(Thread.java:1355)

at

org.apache.hadoop.hdfs.DFSOutputStream\$DataStreamer.closeResponder(DFSOutputStream.java:967)

at

org.apache.hadoop.hdfs.DFSOutputStream\$DataStreamer.endBlock(DFSOutputStream.java:705)

at

org.apache.hadoop.hdfs.DFSOutputStream\$DataStreamer.run(DFSOutputStream.java:894)

25/04/22 01:36:52 WARN hdfs.DFSClient: Caught exception

```
java.lang.InterruptedException
  at java.lang.Object.wait(Native Method)
  at java.lang.Thread.join(Thread.java:1281)
  at java.lang.Thread.join(Thread.java:1355)
  at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStre
am.java:967)
  at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.jav
a:705)
  at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
25/04/22 01:36:52 INFO mapreduce.JobSubmitter: number of splits:1
25/04/22 01:36:53 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1745309395132_0001
25/04/22 01:36:55 INFO impl.YarnClientImpl: Submitted application
application_1745309395132_0001
25/04/22 01:36:55 INFO mapreduce.Job: The url to track the job:
http://quickstart.cloudera:8088/proxy/application_1745309395132_0001/
25/04/22 01:36:55 INFO mapreduce.Job: Running job: job_1745309395132_0001
25/04/22 01:37:16 INFO mapreduce.Job: Job job_1745309395132_0001 running in uber
mode : false
25/04/22 01:37:16 INFO mapreduce.Job: map 0% reduce 0%
25/04/22 01:37:47 INFO mapreduce.Job: map 100% reduce 0%
25/04/22 01:38:05 INFO mapreduce.Job: map 100% reduce 100%
25/04/22 01:38:06 INFO mapreduce.Job: Job job_1745309395132_0001 completed
successfully
25/04/22 01:38:07 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=154
    FILE: Number of bytes written=286987
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=188
    HDFS: Number of bytes written=96
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=27952
    Total time spent by all reduces in occupied slots (ms)=14497
    Total time spent by all map tasks (ms)=27952
    Total time spent by all reduce tasks (ms)=14497
    Total vcore-milliseconds taken by all map tasks=27952
```

Total vcore-milliseconds taken by all reduce tasks=14497
Total megabyte-milliseconds taken by all map tasks=28622848
Total megabyte-milliseconds taken by all reduce tasks=14844928

Map-Reduce Framework

Map input records=6
Map output records=14
Map output bytes=130
Map output materialized bytes=154
Input split bytes=114
Combine input records=14
Combine output records=13
Reduce input groups=13
Reduce shuffle bytes=154
Reduce input records=13
Reduce output records=13
Spilled Records=26
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=401
CPU time spent (ms)=2470
Physical memory (bytes) snapshot=342106112
Virtual memory (bytes) snapshot=3015086080
Total committed heap usage (bytes)=226365440

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=74

File Output Format Counters

Bytes Written=96

```
[cloudera@quickstart ~]$ hdfs dfs -ls /out1
```

Found 2 items

```
-rw-r--r--  1 cloudera supergroup    0 2025-04-22 01:38 /out1/_SUCCESS
-rw-r--r--  1 cloudera supergroup   96 2025-04-22 01:38 /out1/part-r-00000
```

```
[cloudera@quickstart ~]$ hdfs dfs -ls /out1/part-r-00000
```

```
-rw-r--r--  1 cloudera supergroup   96 2025-04-22 01:38 /out1/part-r-00000
```

```
[cloudera@quickstart ~]$ hdfs dfs -cat /out1/part-r-00000
```

bharat 1

boss 1

devansh 1

hp 1

jay 2

ke 1

mata 1
om 1
papadiwal 1
patil 1
ram 1
rohit 1
shree 1

ASSIGNMENT 2

(MAPREDUCER)

REFERENCES :-

<https://mapreduceprogram.blogspot.com/2024/02/mapreduce-program-in-hadoop-by-using.html>

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class MapReduce {

    public static class IPMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable one = new IntWritable(1);
        private Text ip = new Text();

        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
            // Assuming the IP address is the first token in each line
            StringTokenizer itr = new StringTokenizer(value.toString());
            if (itr.hasMoreTokens()) {
                ip.set(itr.nextToken());
                context.write(ip, one);
            }
        }
    }
}
```

```
    }  
    }  
}
```

```
public static class IntSumReducer  
    extends Reducer<Text,IntWritable,Text,IntWritable> {  
    private IntWritable result = new IntWritable();  
  
    public void reduce(Text key, Iterable<IntWritable> values,  
        Context context  
        ) throws IOException, InterruptedException {  
        int sum = 0;  
        for (IntWritable val : values) {  
            sum += val.get();  
        }  
        result.set(sum);  
        context.write(key, result);  
    }  
}
```

```
public static void main(String[] args) throws Exception {  
    Configuration conf = new Configuration();  
    Job job = Job.getInstance(conf, "IP address count");  
    job.setJarByClass(MapReduce.class);  
    job.setMapperClass(IPMapper.class);  
    job.setCombinerClass(IntSumReducer.class);  
    job.setReducerClass(IntSumReducer.class);  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
    FileInputFormat.addInputPath(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
    System.exit(job.waitForCompletion(true) ? 0 : 1);  
}  
}
```